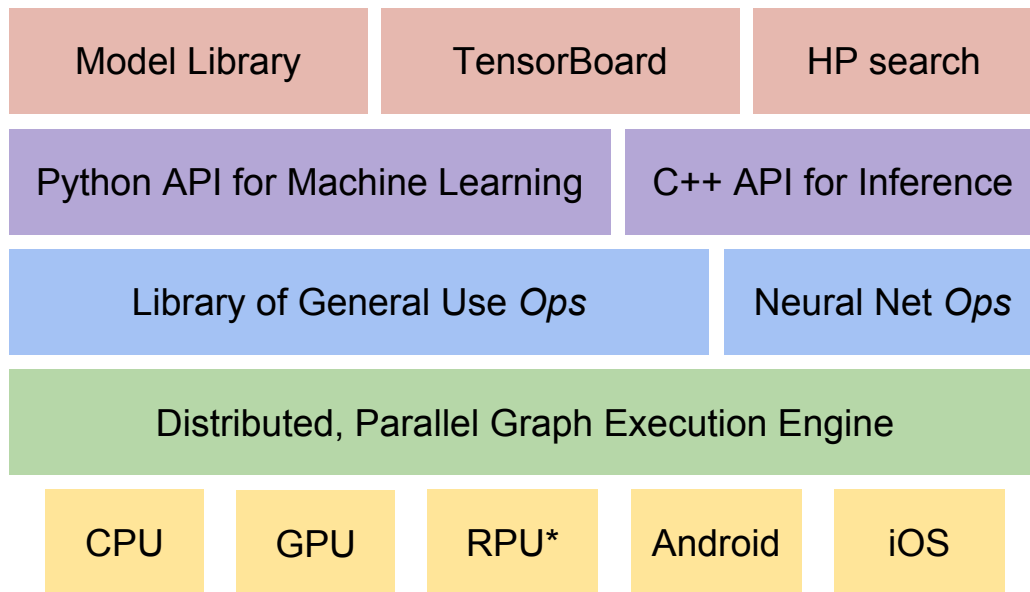


TensorFlow Queues

And Input Pipelines

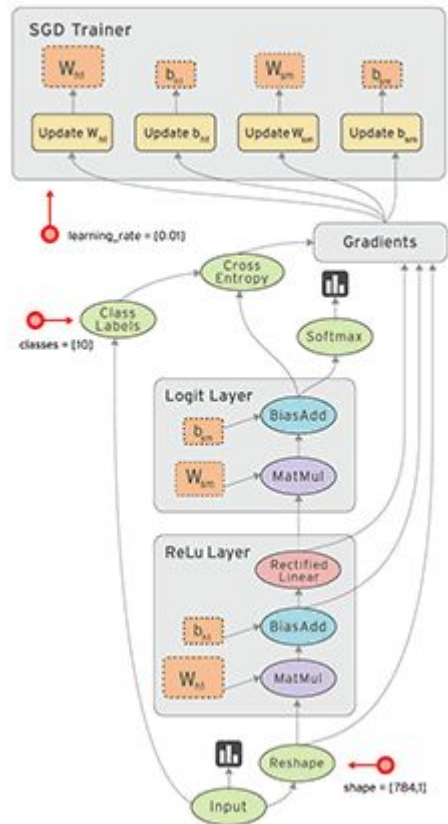
TensorFlow

- System for computation across heterogeneous devices
- Arbitrary operations, arbitrary devices, connected through C interface



Parallel Execution

- Launch graph in a *Session*
- Request output of some Ops with *Run* API
- TensorFlow computes set of Ops that must run to compute the requested outputs
- Ops execute, in parallel, as soon as their inputs are available

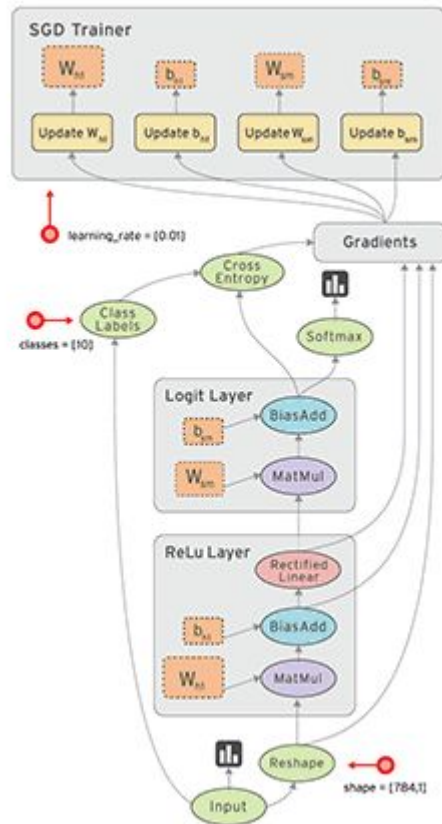


Parallel Execution with Python

1 core

```
for d in data:  
    preprocess(d)  
    Python_code  
    Python_code  
    Python_code
```

48 cores (2x Intel Xeon)
16 GPUs (8 x K80)



GIL!!!!!!



Hacks around GIL

Python-code

Python code
Python code
Numpy.complex_calculation
...
...

Python code

Py_BEGIN_ALLOW_THREADS



Py_END_ALLOW_THREADS

Getting your data into TensorFlow

1. `tf.constant(mydata)` BAD
2. `session.run(node, feed_dict={in:data})` BETTER
3. Queues and Input Pipelines BEST!

tf.constant(data)

- Used in [tensorflow/examples/how_tos/reading_data/fully_connected_preloaded.py](#)

```
with tf.device('/cpu:0'):
    input_images = tf.constant(data_sets.train.images)
    input_labels = tf.constant(data_sets.train.labels)
```

- Inlines data in the Graph definition
- hard 2GB limit for size of Graph
- Single-threaded (Graph not thread-safe)
- Running remotely adds additional protobuf encoding/decoding step

`sess.run(..., feed_dict=...)`

- Gets pointer to underlying numpy memory buffer
- Single-threaded memcopy (<2GB/s max)

BUT:

- Could be replaced with multi-threaded memcopy in future
- Can do `sess.run` calls in parallel, `sess.run` releases GIL

`sess.run(..., feed_dict=...)`

- What if you need to preprocess the data?

```
im = random_crop(im)
im = random_flip_left_right(im)
im = random_brightness(im)
im = random_contrast(im)
im = per_image_whitening(im)
```

- What if you need to implement pre-fetching/buffering?
 - ...

`sess.run(..., feed_dict=...)`

- What if you need to preprocess the data?

```
im = random_crop(im)
im = random_flip_left_right(im)
im = random_brightness(im)
im = random_contrast(im)
im = per_image_whitening(im)
```



TensorFlow native
implementations

- What if you need to implement pre-fetching/buffering?

-




Queues

Queues


`a=sess.run`

Enqueue



Buffer

Dequeue



`b=sess.run`

Queues


`a=sess.run`

Enqueue

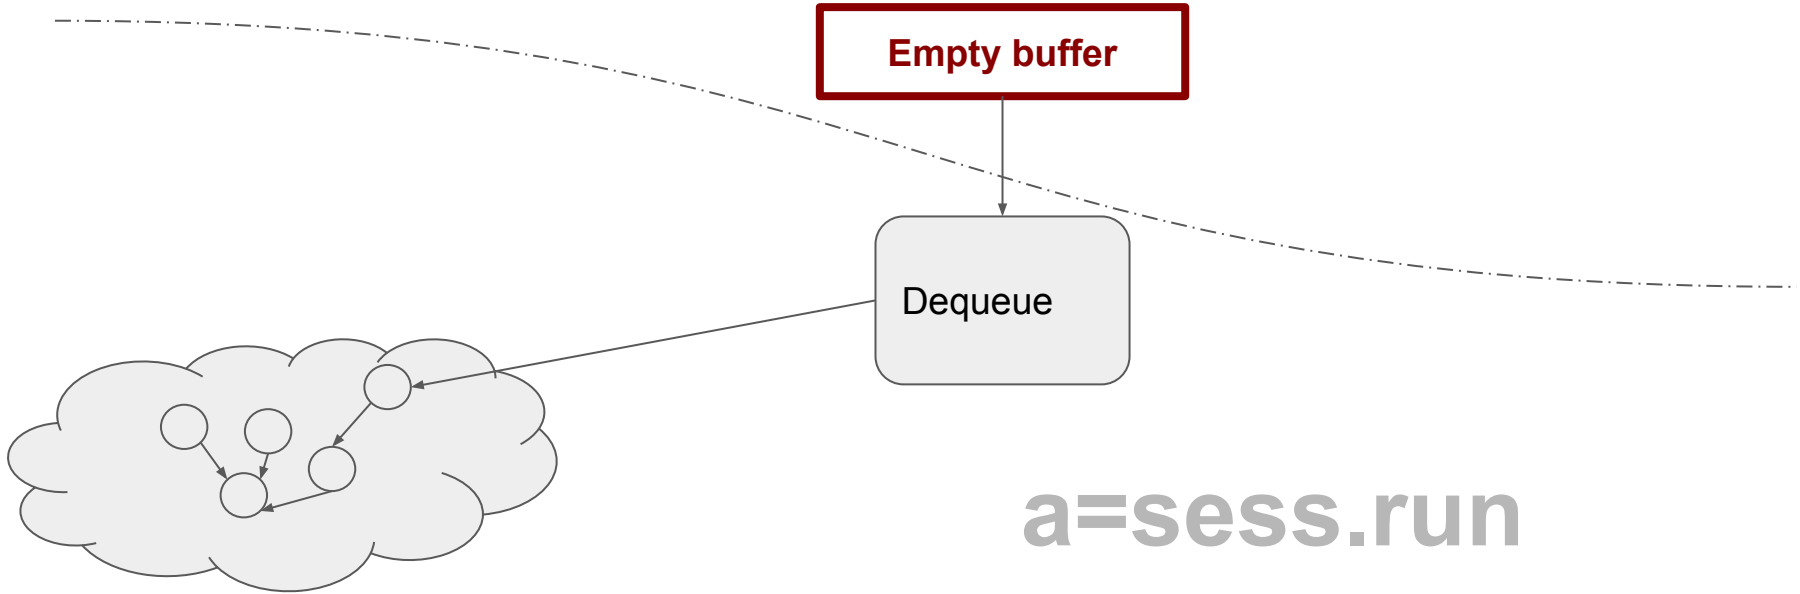


Full Buffer

**Hang
forever**



Hang forever



Hang forever unless

Session was configured with timeouts

```
session = tf.Session(tf.Config(operation_timeout_in_ms=2000))
```

DeadlineExceededError

OR

Queue was closed by pushing a special “Close” token on it.

```
sess.run(queue.close())
```

OutOfRangeError/AbortedError

Queue Example

```
q = tf.FIFOQueue(capacity=20, dtypes=[tf.int32])
enqueue_placeholder = tf.placeholder(tf.int32)
enqueue_op = q.enqueue(enqueue_placeholder)
sess = create_session()
for i in range(10):
    sess.run(enqueue_op, feed_dict={enqueue_placeholder:i})
    print "Queue size is now: "+str(sess.run(q.size()))
sess.run(q.close())
```


Queue Example

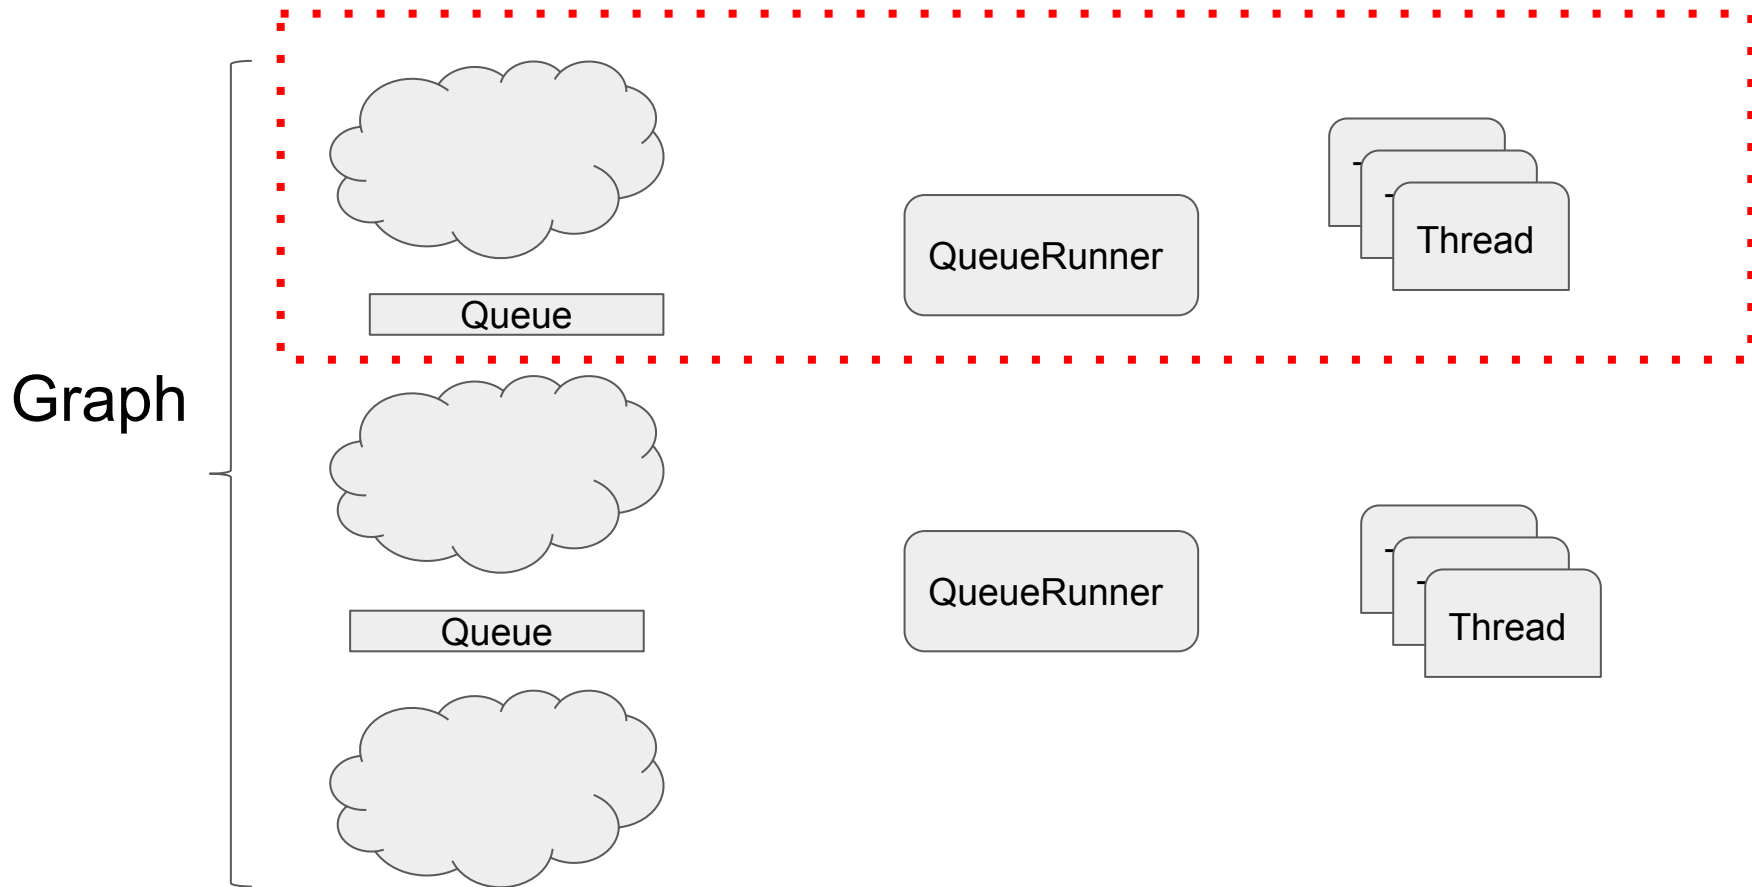
```
random_number = tf.random_uniform(shape=())  
q = tf.FIFOQueue(capacity=20, dtypes=[tf.float32])  
enqueue_op = q.enqueue(random_number)
```

```
sess = create_session()  
print sess.run(q.size()) # prints 0  
def run():  
    for i in range(5):  
        sess.run(enqueue_op)
```

```
threads = [threading.Thread(target=run, args=()) for i in range(2)]  
[t.start() for t in threads]  
print sess.run(q.size()) # prints 4  
time.sleep(0.5)  
print sess.run(q.size()) # prints 10
```

QueueRunners

InputProducer



Gotchas

- Queues are stateful

```
merged = make_image_input()
train_step = make_train_step(merged)
for step in range(5):
    loss_value = sess.run(train_step, feed_dict={y_label:y_1})
    summary_str = sess.run(merged, feed_dict={y_label:y_1})
    writer.add_summary(summary_str, step)
```

where did extra images go?

Gotchas

```
ranges = tf.train.range_input_producer()  
number = ranges.dequeue()  
sess = create_session  
tf.train.start_queue_runners()
```

```
if version==1:  
    print sess.run([number]*3)  
elif version ==2:  
    print sess.run([ranges.dequeue()]*3)  
elif version == 3:  
    print sess.run([ranges.dequeue(), ranges.dequeue(), ranges.dequeue()])
```