

Memory

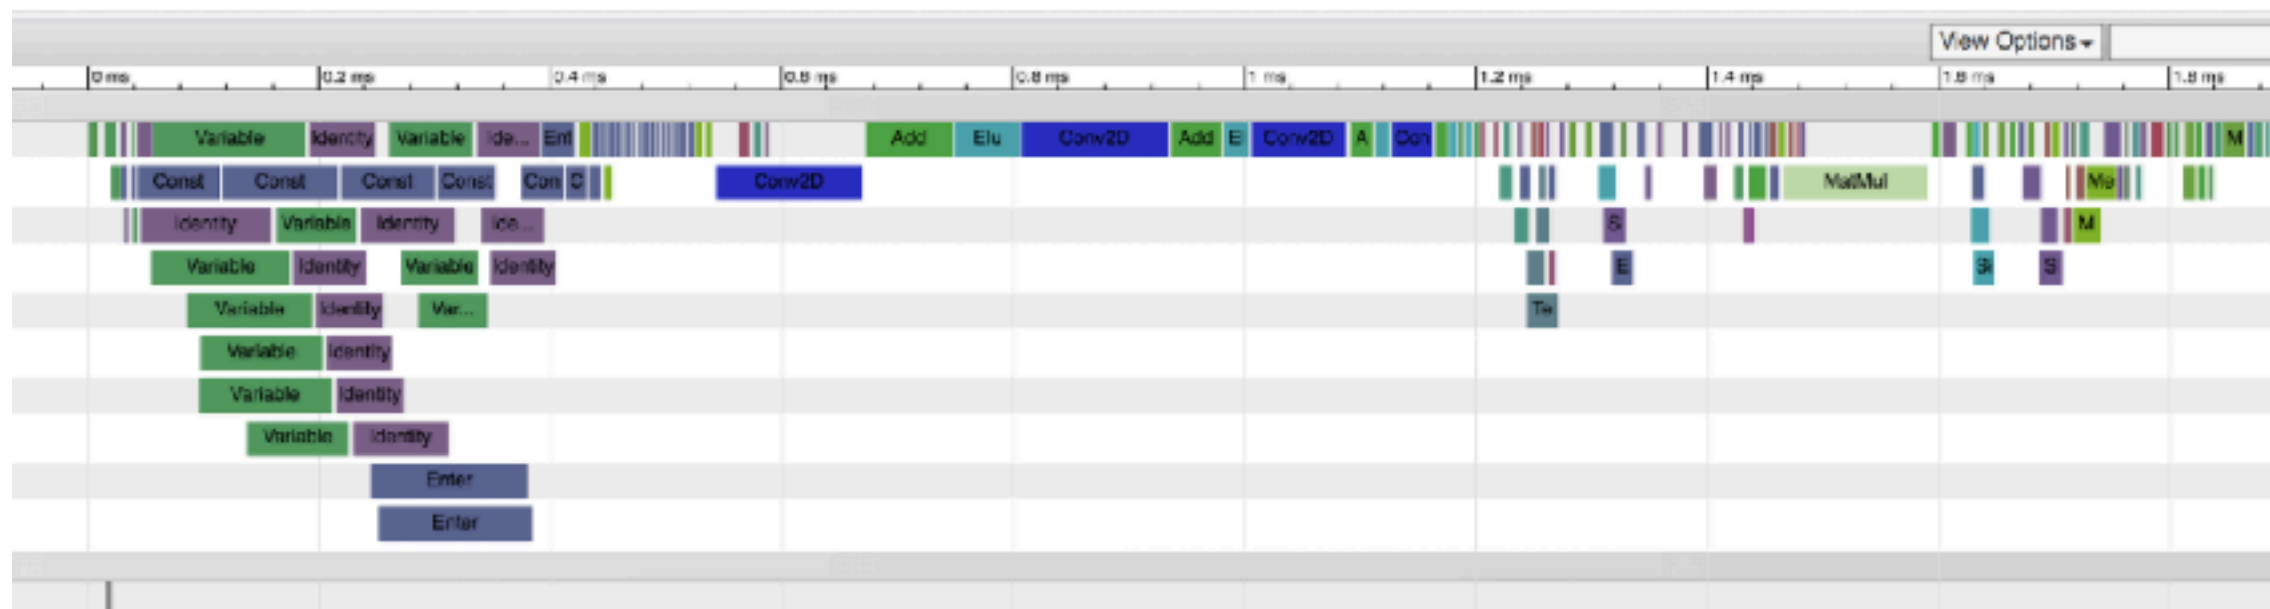
in neural networks (TensorFlow)

Why memory

- Peak efficiency is achieved for large batches on GPU:
 - 11 T ops/s for 8k-by-8k matrix multiply on TitanX
 - 1.1 T ops/sec for 8k-by-8k matrix multiply on Xeon V3
 - 0.1 G ops/s for 256-by-256 matrix multiply on TitanX/Xeon
 - Small batch runtime dominated by once-per-batch overhead (ie, var reads)
- Money doesn't buy more GPU memory
 - \$800 — 1080TI: 12GB
 - \$6k — P100: 16GB
 - \$??? — V100: 16GB


Fprop efficiency

Universe starter agent: 400 fps...too slow




Fprop efficiency

1	565
2	1089
3	1691
4	2196
5	2712
6	3320
7	3569
8	4006
16	8275
32	12664
64	15840
128	22840
256	22415




Fprop efficiency




Calculating memory requirements

- $f(x) + g(x)$
- 2 memory units










Calculating memory requirements

- $f(x) + g(x) + h(x)$
- either 2 or 3 memory units



Pebble game


- Rules:
 1. can only put pebble on node if all parents have pebbles
 2. goal to put pebble on final node



- Rules:

1. can only put pebble on node if all parents have pebbles or no parents
2. goal to put pebble on final node

Example



Best case: 3 units

Worst case: N units


TensorFlow case: ???

Pebble game



- Number of pebbles needed = peak requirement
 - Different schedules produce different requirements
 - How to find the optimal schedule?
 - No solution for general computation graphs
- “Inapproximability of treewidth, one-shot pebbling, and related layout problems” <http://dl.acm.org/citation.cfm?id=2655729>*
- But good heuristics exist

Pebble game

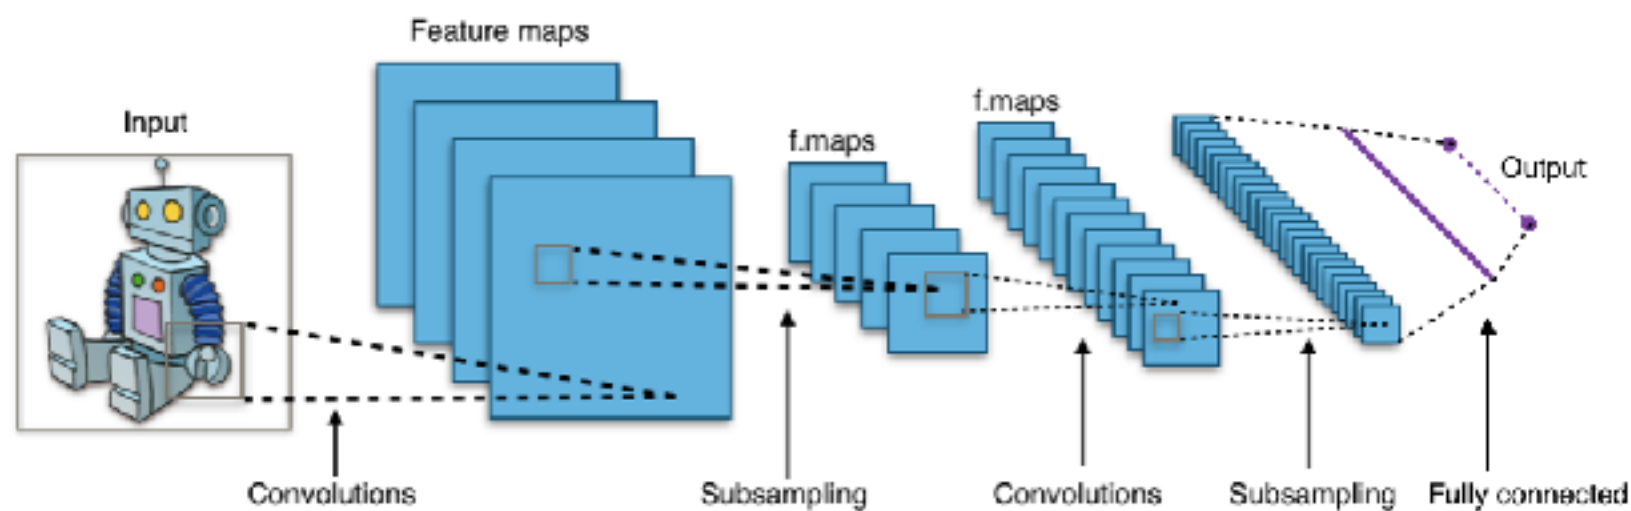

- One-shot pebbling = do not touch nodes already visited = no recomputations
- Multi-shot pebbling = can revisit old nodes = recomputations allowed
- TensorFlow = no recomputations



Neural networks



Inference



**Memory requirement determined by most expensive layer
(typically the first fully connected layer)**

Training


$$c = f(g(h(x)))$$

$$\frac{dc}{dx} = f'(g(h(x)))g'(h(x))h'(x)$$

a3
a2
a1

$$\frac{dc}{dx} = f'(g(h(x)))g'(h(x))h'(x)$$

b1
b2
b3



Training

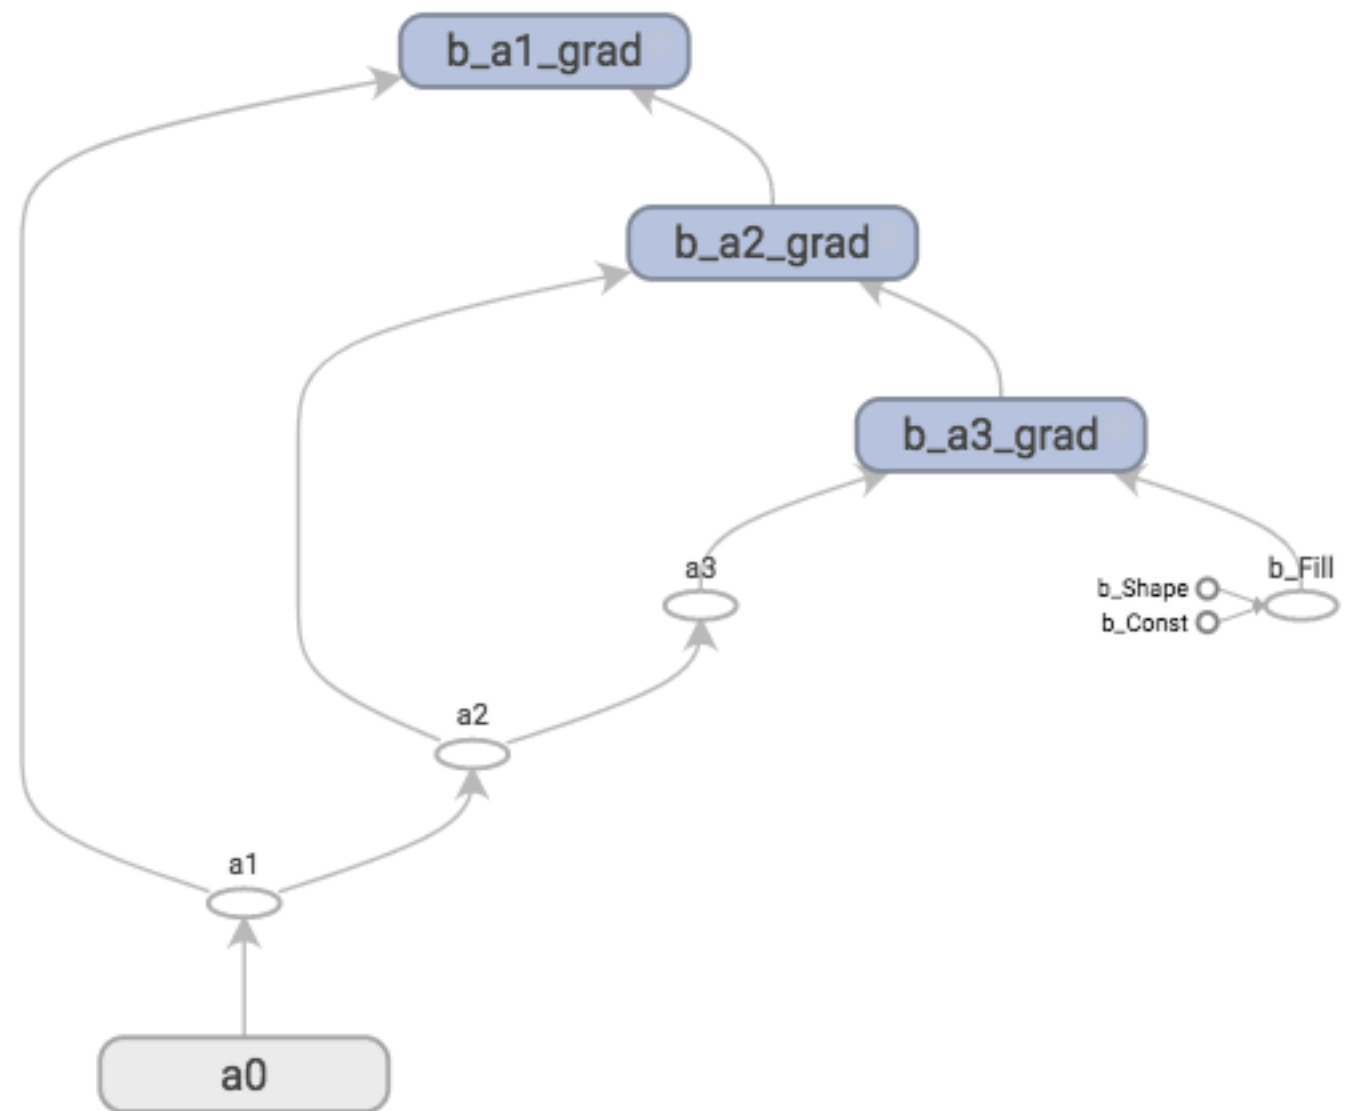
```
3]: tf.reset_default_graph()
node_mbs = 1
length = 4

dtype = np.float32
n = node_mbs * 250000
a0_ = tf.ones((n,), dtype=dtype)
a0 = tf.Variable(a0_, name="a0")
a = a0
for i in range(1, length):
    name = "a"+str(i)
    a = tf.tanh(a, name=name)

grad = tf.gradients([a], [a0])[0]
sess = create_session()
```

```
4]: show_graph(ungroup_gradients=True)
```

Main Graph



TensorFlow memory


- Which order does it pick? (look in [executor.cc](#))

TensorFlow memory




```
result = tf.random_uniform((size, size))
```

```
for i in range(n):
```

```
    result = result @ tf.random_uniform((size, size))
```



TensorFlow memory



8k matmul takes 100ms

8k-by-8k random_uniform takes 4ms

How to monitor memory

- TensorFlow manages its own memory, so nvidia-smi is useless

1. parse LOG_MEMORY allocation/deallocation messages
(https://github.com/yaroslavvb/memory_util)

2. Extract it from Timeline

3. Write custom TensorFlow op that queries allocator on demand

https://github.com/yaroslavvb/memory_probe_ops

memory_util example

- <https://github.com/yaroslavvb/notebooks/blob/master/mnist-memory.ipynb>
- https://github.com/yaroslavvb/memory_util

timeline

```
run_metadata = tf.RunMetadata()  
run_options = tf.RunOptions(trace_level=tf.RunOptions.FULL_TRACE)  
sess.run(model.train_op, options=run_options, run_metadata=run_metadata)
```

```
node_name: "a02_add"  
all_start_micros: 1505768360742529  
op_start_rel_micros: 32  
op_end_rel_micros: 80  
all_end_rel_micros: 137  
memory {  
  allocator_name: "GPU_0_bfc"  
  allocator_bytes_in_use: 171016448  
}  
output {  
  tensor_description {  
    dtype: DT_FLOAT  
    shape {  
      dim {  
        size: 250000  
      }  
    }  
    allocation_description {  
      requested_bytes: 1000000  
      allocated_bytes: 1000192  
      allocator_name: "GPU_0_bfc"  
      allocation_id: 3735  
      ptr: 1106455317760  
    }  
  }  
}  
timeline_label: "a02_add = Add[a01_add, a02_tanh]"  
scheduled_micros: 1505768360742492  
memory_stats {  
}
```


<https://github.com/yaroslavvb/stuff/blob/master/memory%20tracking.ipynb>

Improving memory usage

- Pick better execution order
- Forget/recompute intermediate Tensors
- Use TensorFlow functions
- Offload to main memory

Better execution order

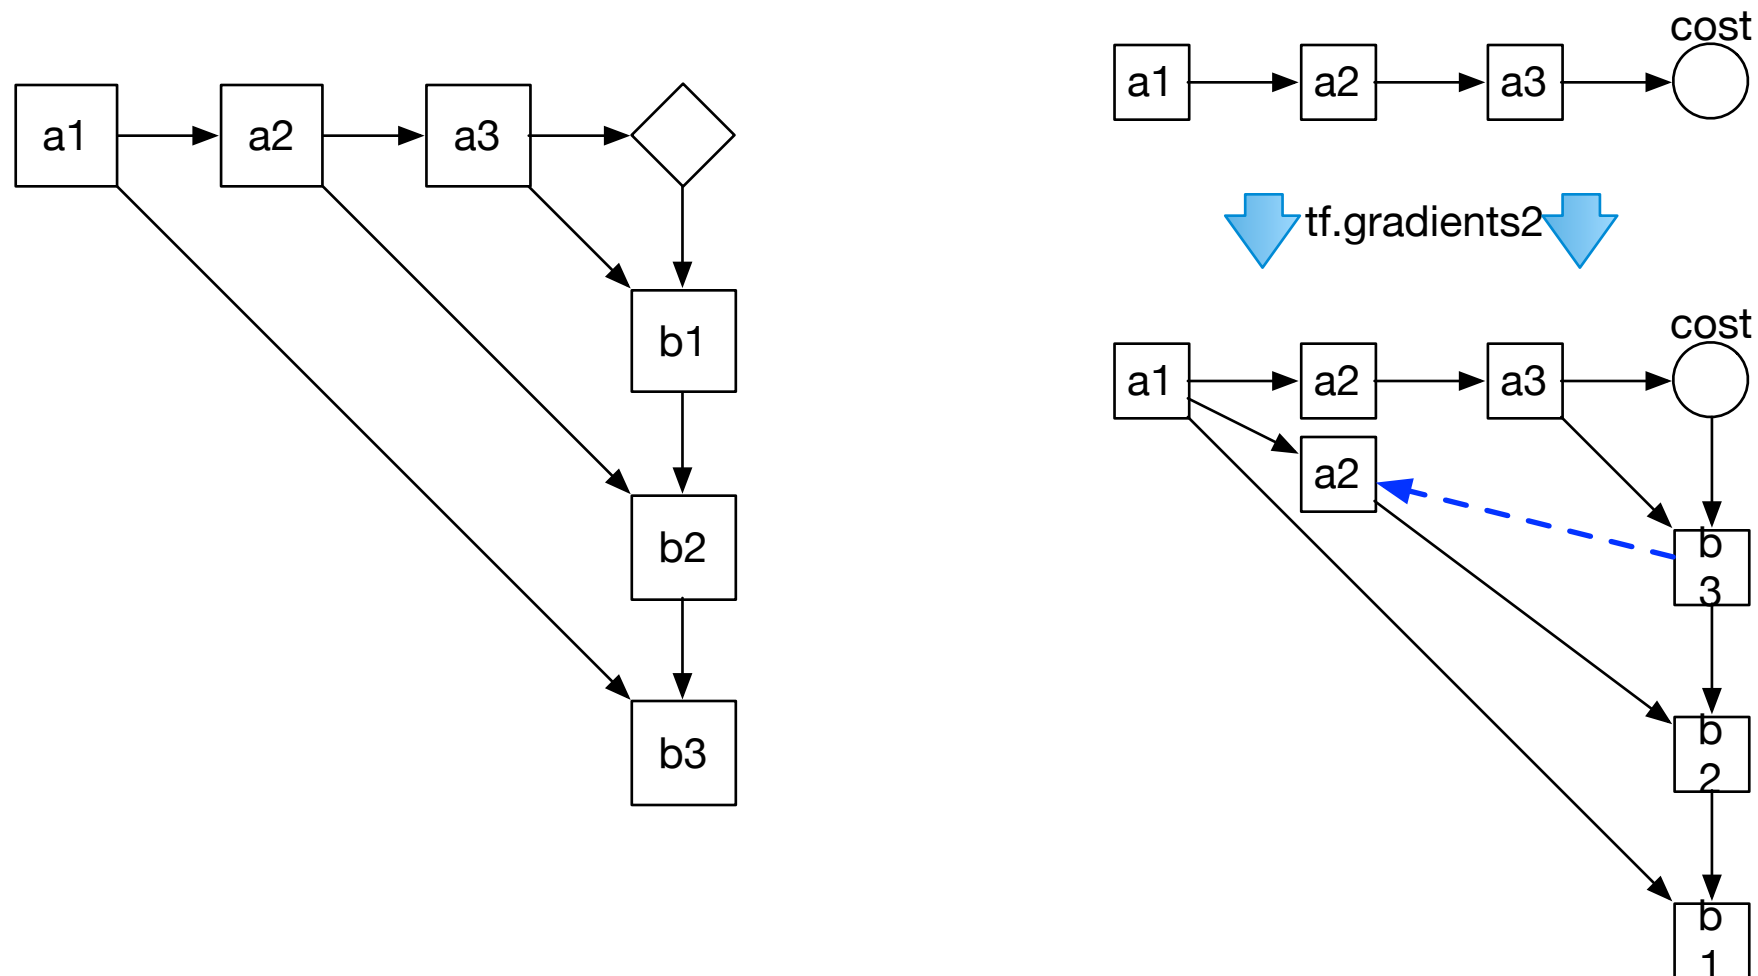
Add enough control dependencies so that execution order is deterministic.




Pick execution order where nodes that are needed later, are also computed later

<https://github.com/yaroslavvb/stuff/tree/master/linearize>

Rewire the graph for recompilation




$\text{sqrt}(n)$ saving



Training Deep Nets with Sublinear Memory Cost

[Tianqi Chen](#), [Bing Xu](#), [Chiyuan Zhang](#), [Carlos Guestrin](#)





Functions to recompute intermediate results

```
@function.Defun(tf.float32, func_name="tanh3")  
def tanh3(a):  
    return tf.tanh(tf.tanh(tf.tanh(a)))
```

intermediate results are forgotten

similar to graph rewiring, but requires modifying model construction code
technique used by Google Translation and LM models



<https://github.com/yaroslavvb/stuff/blob/master/saving%20memory%20by%20using%20functions.ipynb>

Offload to main RAM

- Instead of forgetting/recomputing, save to main memory (rewriting graph, using `swap_memory=True`, or grappler)
- makes sense for $O(n^3)$ ops (conv2d, matmul)
- doesn't make sense for $O(n^2)$ ops (everything else)
- 7x faster to recompute `tf.mul` on GPU than load from memory (10x faster for `tf.concat`)
- <https://github.com/yaroslavvb/stuff/blob/master/gpu-memory-transfer.ipynb>